

Home Page - Select or creat SOFT_NN - Jupyter Not +

localhost:8888/notebooks/SOFT_NN.ipynb

jupyter SOFT_NN Last Checkpoint 18 hours ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3

```
In [1]: import os
import keras
using TensorFlow backend.

In [2]: from __future__ import absolute_import
from __future__ import division
from __future__ import print_function

import itertools

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from pylab import rcParams
import matplotlib

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
train = pd.read_csv(r'C:\Users\Withishma\Desktop\train.csv')
print('Shape of the train data with all features:', train.shape)
train = train.select_dtypes(exclude=['object'])
print("")
print('Shape of the train data with numerical features:', train.shape)
train.drop('Id',axis = 1, inplace = True)
train.fillna(0,inplace=True)
test = pd.read_csv(r'C:\Users\Withishma\Desktop\test.csv')
test = test.select_dtypes(exclude=['object'])
ID = test.Id
test.fillna(0,inplace=True)
test.drop('Id',axis = 1, inplace = True)

print("")
print("List of features contained our dataset:",list(train.columns))

Shape of the train data with all features: (1460, 81)

Shape of the train data with numerical features: (1460, 38)

List of features contained our dataset: ['MSSubClass', 'LotFrontage', 'LotArea', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd', 'MasVnrArea', 'BsmtFinS1', 'BsmtFinS2', 'BsmtUnfSF', 'TotalBsmtSF', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'TotRmsAbvGrd', 'Fireplaces', 'GarageVeh', 'GarageCars', 'GarageArea', 'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch', '3SeasonPorch', 'ScreenPorch', 'PoolArea', 'MiscVal', 'MSold', 'YrSold', 'SalePrice']
```

Type here to search

ENG 18:14 03-11-2019

Home Page - Select or creat SOFT_NN - Jupyter Not +

localhost8888/notebooks/SOFT_NN.ipynb

Jupyter SOFT_NN Last Checkpoint: 18 hours ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3

```
In [3]: from sklearn.ensemble import IsolationForest

clf = IsolationForest(max_samples = 100, random_state = 42)
clf.fit(train)
y_noano = clf.predict(train)
y_noano = pd.DataFrame(y_noano, columns = ['Top'])
y_noano[y_noano['Top'] == 1].index.values

train = train.loc[y_noano['Top'] == 1].index.values
train.reset_index(drop = True, inplace = True)
print("Number of outliers:", y_noano[y_noano['Top'] == -1].shape[0])
print("Number of rows without outliers:", train.shape[0])

C:\Users\Withishma\Anaconda3\lib\site-packages\sklearn\ensemble\isoforest.py:237: FutureWarning: default contamination parameter 0.1 will change in version 0.22 to "auto". This will change the predict method behavior.
FutureWarning)
C:\Users\Withishma\Anaconda3\lib\site-packages\sklearn\ensemble\isoforest.py:247: FutureWarning: behaviour="old" is deprecated and d will be removed in version 0.22. Please use behaviour="new", which makes the decision_function change to match other anomaly detection algorithm API.
FutureWarning)

Number of outliers: 146
Number of rows without outliers: 1314

C:\Users\Withishma\Anaconda3\lib\site-packages\sklearn\ensemble\isoforest.py:415: DeprecationWarning: threshold_attribute is deprecated in 0.20 and will be removed in 0.22.
" removed in 0.22.". DeprecationWarning)

In [4]: train.head(10)

Out[4]:
```

	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea	BsmfFinSF1	BsmfFinSF2	WoodDeckSF	OpenPor
0	80	65.0	8450	7	5	2003	2003	196.0	706	0	...	0
1	20	80.0	9600	6	8	1976	1976	0.0	678	0	...	268
2	60	68.0	11250	7	5	2001	2002	162.0	486	0	...	0
3	70	60.0	9550	7	5	1915	1970	0.0	218	0	...	0
4	60	84.0	14200	8	5	2000	2000	350.0	655	0	...	192
5	50	85.0	14115	5	5	1993	1995	0.0	732	0	...	40
6	20	75.0	10084	8	5	2004	2005	186.0	1369	0	...	255
7	80	61.0	6120	7	5	1931	1950	0.0	0	0	...	90

Type here to search

Home Page - Select or creat SOFT_NN - Jupyter Not +

localhost8888/notebooks/SOFT_NN.ipynb

Jupyter SOFT_NN Last Checkpoint: 18 hours ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3

```
In [5]: import warnings
warnings.filterwarnings('ignore')

col_train = list(train.columns)
col_train_bis = list(train.columns)

col_train_bis.remove('SalePrice')

mat_train = np.matrix(train)
mat_test = np.matrix(test)
mat_new = np.matrix(train.drop('SalePrice',axis = 1))
mat_y = np.array(train.SalePrice).reshape((1314,1))

prepro_y = MinMaxScaler()
prepro_y.fit(mat_y)

prepro = MinMaxScaler()
prepro.fit(mat_train)

prepro_test = MinMaxScaler()
prepro_test.fit(mat_new)

train = pd.DataFrame(prepro.transform(mat_train),columns = col_train)
test = pd.DataFrame(prepro_test.transform(mat_test),columns = col_train_bis)
train.head()

Out[5]:
```

	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea	BsmfFinSF1	BsmfFinSF2	WoodDeckSF	OpenPor
0	0.235294	0.207668	0.028202	0.625	0.428571	0.953488	0.883333	0.12250	0.416274	0.0	...	0.000000
1	0.000000	0.255591	0.072904	0.500	0.857143	0.744196	0.433333	0.00000	0.576951	0.0	...	0.404891
2	0.235294	0.217262	0.087396	0.625	0.428571	0.937694	0.696967	0.10125	0.286557	0.0	...	0.000000
3	0.294118	0.191963	0.072464	0.625	0.428571	0.271318	0.333333	0.00000	0.127358	0.0	...	0.000000
4	0.235294	0.268371	0.113635	0.750	0.428571	0.930233	0.833333	0.21875	0.386203	0.0	...	0.260870

5 rows x 37 columns

```
In [6]: COLUMNS = col_train
FEATURES = col_train_bis
```

Type here to search

ENG IN 03-11-2019

Home Page - Select or creat SOFT_NN - Jupyter Not X +

localhost8888/notebooks/SOFT_NN.ipynb

Jupyter SOFT_NN Last Checkpoint: 18 hours ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3

```
In [6]: COLUMNS = col_train
FEATURES = col_train_bis
LABEL = "SalePrice"

# Columns
feature_cols = FEATURES

# Training set and Prediction set with the features to predict
training_set = train[COLUMNS]
prediction_set = train.SalePrice

# Train and Test
x_train, x_test, y_train, y_test = train_test_split(training_set[FEATURES], prediction_set, test_size=0.33, random_state=42)
y_train = pd.DataFrame(y_train, columns = [LABEL])
training_set = pd.DataFrame(x_train, columns = FEATURES).merge(y_train, left_index = True, right_index = True)
training_set.head()

# Training for submission
training_sub = training_set[col_train]

In [7]: y_test = pd.DataFrame(y_test, columns = [LABEL])
testing_set = pd.DataFrame(x_test, columns = FEATURES).merge(y_test, left_index = True, right_index = True)
testing_set.head()

Out[7]:
```

	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea	BsmFtnSF1	BsmFtnSF2	...	WoodDeckSF	Open
1232	0.235294	0.239817	0.074221	0.825	0.571429	0.614729	0.800000	0.000000	0.574882	0.0	...	0.343750	
76	0.352941	0.000000	0.083505	0.825	0.428571	0.891473	0.786887	0.01375	0.000000	0.0	...	0.183043	
962	0.235294	0.249201	0.091349	0.500	0.571429	0.682171	0.300000	0.000000	0.175708	0.0	...	0.000000	
433	0.284118	0.159744	0.010540	0.825	0.857143	0.271318	0.918987	0.000000	0.178297	0.0	...	0.000000	
1110	0.204118	0.308709	0.103927	0.375	0.428571	0.283566	0.000000	0.000000	0.000000	0.0	...	0.275815	

5 rows x 37 columns

```
In [8]: import numpy as np
from keras.models import Sequential
from keras.layers import Dense
from keras.wrappers.scikit_learn import KerasRegressor

seed = 7
np.random.seed(seed)

# Model
model = Sequential()
model.add(Dense(100, input_dim=36, kernel_initializer='normal', activation='relu'))
model.add(Dense(100, kernel_initializer='normal', activation='relu'))
model.add(Dense(50, kernel_initializer='normal', activation='relu'))
model.add(Dense(25, kernel_initializer='normal', activation='relu'))
model.add(Dense(1, kernel_initializer='normal'))
# compile model
model.compile(loss='mean_squared_error', optimizer=keras.optimizers.Adam())

feature_cols = training_set[FEATURES]
labels = training_set[LABEL].values

model.fit(np.array(feature_cols), np.array(labels), epochs=100, batch_size=10)

Epoch 80/100 [=====] - 0s 281us/step - loss: 0.0017
Epoch 88/100 [=====] - 0s 355us/step - loss: 0.0023
Epoch 89/100 [=====] - 0s 306us/step - loss: 0.0018
Epoch 90/100 [=====] - 0s 297us/step - loss: 0.0019
Epoch 91/100 [=====] - 0s 258us/step - loss: 0.0018
Epoch 92/100 [=====] - 0s 335us/step - loss: 0.0019
Epoch 93/100 [=====] - 0s 328us/step - loss: 0.0017
Epoch 94/100 [=====] - 0s 338us/step - loss: 0.0019
Epoch 95/100 [=====] - 0s 283us/step - loss: 0.0019
Epoch 96/100 [=====] - 0s 390us/step - loss: 0.0016

In [9]: # Evaluation on the test set created by train_test_split
model.evaluate(x_test, y_test, batch_size=10)
```

Type here to search

ENG IN 18:14 03-11-2019

Jupyter Notebook interface showing code execution for a machine learning model. The code includes training, evaluation, and prediction steps. The output shows the model's performance metrics and the predicted values for the test set.

```
Epoch 95/100
880/880 [=====] - 0s 283us/step - loss: 0.0019
Epoch 96/100
880/880 [=====] - 0s 396us/step - loss: 0.0016
-----
In [9]: # Evaluation on the test set created by train_test_split
model.evaluate(np.array(feature_cols), np.array(labels))
880/880 [=====] - 0s 284us/step
Out[9]: 0.0016664503078738396

In [10]: feature_cols_test = testing_set[FEATURES]
labels_test = testing_set[LABEL].values
y = model.predict(np.array(feature_cols_test))
predictions = list(itertools.islice(y, testing_set.shape[0]))

In [11]: predictions = prepro_y.inverse_transform(np.array(predictions).reshape(434,1))

In [13]: reality = pd.DataFrame(prepro.inverse_transform(testing_set), columns = [COLUMNS].SalePrice

In [14]: y_predict = model.predict(np.array(test))

def to_submit(pred_y_name_out):
    y_predict = list(itertools.islice(pred_y, test.shape[0]))
    y_predict = pd.DataFrame(prepro_y.inverse_transform(np.array(y_predict).reshape(len(y_predict),1)), columns = ['SalePrice'])
    y_predict = y_predict.join(ID)
    y_predict.to_csv(name_out + '.csv', index=False)
    to_submit(y_predict, "Realestate_prices")

In [ ]:
```

Below the Jupyter Notebook, a Microsoft Excel spreadsheet titled "Realestate_prices.csv" is displayed. The spreadsheet contains two columns: "SalePrice" and "Id". The data is as follows:

SalePrice	Id
125377.46	1461
139463.03	1462
187838.8	1463
197409.16	1464
200685.45	1465
187992.8	1466
184884.81	1467
177838	1468
196438.4	1469
123465.02	1470
217517.52	1471
103710.68	1472
104635.375	1473
149265.44	1474
114954.9	1475
365955.9	1476
260562.39	1477
303502.47	1478
288211.72	1479
488054.6	1480
354869.62	1481
213429.3	1482
184882.4	1483
177184.34	1484